

**CSL373/CSL633 Minor 1 Exam**  
**Operating Systems**  
**Sem II, 2013-14**

Answer all 6 questions

Max. Marks: 32

**Unix System Calls**

1. What is the total number of processes at the end of the execution of the following program? Assume there is one process in the beginning that starts running at main. Also, assume that all system calls succeed.

```
main() {  
    fork();  
    fork();  
    fork();  
}
```

Explain. [3]

Ans:  $2^3 = 8$ .

fork#	#processes
1	2
2	4
3	8

First fork is called by the first process, it generates one more process (or doubles). Both these processes call second fork and doubles the number of processes to 4...and so on.

2. Consider the following program:

```
main() {  
    int fd;  
    fd = open("outfile", O_RDWR)  
    fork();  
    write(fd, "hello", 5);  
    exit();  
}
```

Assume all system calls finish successfully on a uniprocessor system. Also, assume that a system call cannot be interrupted in the middle of its execution. What will be the contents of the "outfile" file, after all processes have successfully exited? Explain briefly. [3]

hellohello

Fork will create two processes with the same *fd* pointing to the given file.

Two "hello" writes will be observed by the same *fd* (seek value or pos will be shared). And because of the given assumption, interleaving/mixing of these two writes is not possible.

3. Now, consider the following program:

```
main() {  
    int fd;  
  
    fork();  
    fd = open("outfile", O_RDWR);  
    write(fd, "hello", 5);  
    exit();  
}
```

Assume all system calls finish successfully on a uniprocessor system. Also, assume that a system call cannot be interrupted in the middle of its execution.

Notice that open is now called after fork (not before fork as in the previous question). What will be the contents of the "outfile" file, after all processes have successfully exited? Explain briefly. [3]

hello

Value of fd will be same for both the process.

But open file entry corresponding to both this fd are different. ie. there are two copies of write offset.

After open system call, write offset in each of the opened file entry is initialized to zero.

And second write will over write the first one.

## Virtual Memory

4. Consider the following program header of an ELF executable file a.out:

```
LOAD: offset 0x00001000
      vaddr 0x40100000
      paddr 0x00100000//removed this line
      align 2**12
      filesize 0x0000b596
      memsize 0x000126fc
      flags rwx
```

4a. Assume that this executable is loaded using the `exec("a.out", ...)` system call on 32-bit Linux. Also, assume that the Linux kernel is mapped starting at virtual address `0xc0000000`. Draw the layout of the virtual address space of the process just after successful completion of the `exec()` system call. Indicate the sizes, and the contents of the memory regions, wherever possible. [6]

Answer:

```
0xffffffff
...      Kernel (Access: kernel only)
0xc0000000
...      Heap(why?)
...
...      Heap(why?)
0x401126fc
...      zero filled memory (user)
0x4010b596
...      Program (user)
0x40100000
...      (Either unmapped: or Devices/VGA/Bios mapped but access to kernel only).
...
...      Unmapped (atleast will help in catching null pointer dereference at kernel level) or
access to kernel only.
0x00000000
```

4b. Assume that the operating system is using paging to map the pages of the executable on



5. Consider the code in bootasm.S (for the bootsector). At line 8468, the code uses the “call” instruction to branch into the C function called “bootmain” (defined at line 8517 in bootmain.c). Would it have been okay to instead use the “jmp” instruction like this:

```
8467  movl  $start, %esp
8468  jmp   bootmain
```

If this is okay, explain why. If not, explain why not? [3]

Yes: bootmain shouldn't return. So it is okay to assume bootmain as noreturn function and optimize out call statement by replacing with jmp

No, It makes it difficult to debug if there exists a bug: bootmain returns.

6. Consider lines 1049-1051 that enable paging in the hardware. As soon as paging is enabled (at line 1051), the address space changes. i.e., some addresses that were valid previously are no longer valid. Similarly, some (virtual) addresses that were invalid previously are now valid. Assume that the size of the physical memory is 64MB.

```
pde_t entrypgdir[NPDENTRIES] = {
    // Map VA's [0, 4MB) to PA's [0, 4MB)
    [0] = (0) | PTE_P | PTE_W | PTE_PS,
    // Map VA's [KERNBASE, KERNBASE+4MB) to PA's [0, 4MB)
    [KERNBASE>>PDXSHIFT] = (0) | PTE_P | PTE_W | PTE_PS,
};
```

KERNBASE is 0x80000000.

$V2P(X) = X - \text{KERNBASE}$ .

Address range	Before paging	After paging	Example symbols
[0,4MB)	VALID	VALID	_start
[4MB,physmem)	VALID	INVALID	
[KERNBASE, KERNBASE+4MB)	INVALID	VALID	entry,stack,main,entrypgdir

a. Give an example of a (virtual) address that was invalid before enabling paging at line 1051 (i.e., accessing it would have generated an exception), but would be valid after enabling paging. [2]

Ans: Any location between [KERNBASE, KERNBASE+4MB)

Wrong answer: all other values are wrong.

Example symbols located between [KERNBASE,KERNBASE+4MB): entry,stack,main,entrypgdir

Wrong symbols: \_start

Why you will get exceptions before: Accessing value outside physical memory.

Why you won't get exceptions after: Page table mapping exists.

b. Give an example of a (virtual) address that was valid before enabling paging at line 1051 (i.e., accessing it would have resulted in some valid data), but would be invalid after enabling paging i.e., accessing that address will generate an exception). [2]

Ans: Any valid physical address other than [0,4MB) and [KERNBASE,KERNBASE+4MB)

Wrong ans: symbols between [KERNBASE,KERNBASE+4MB): entry,stack,main

Wrong ans: symbols between [0,4MB): \_start

Why you won't get exceptions before:

valid physical memory location (assuming physical memory is more than 4MB)

Why you will get exceptions after:

Only [0,4MB) and [KERNBASE,KERNBASE+4MB) are valid.